Illustrative Enterprise AI Scope of Work

Craft Brewing, Manufacturing & Distribution

(Sample Only — For Demonstration of Approach, Not a Proposal)

1.0 Project Overview

Modern craft breweries operate as hybrid businesses spanning **manufacturing, distribution, hospitality, and brand marketing**. This creates complex operational and data challenges that are difficult to manage with traditional reporting and disconnected systems.

Based on experience deploying AI-enabled systems in regulated, margin-sensitive industries, this document outlines **illustrative enterprise AI applications** that may be appropriate for a regional brewing company *once human workflows and operational constraints are clearly understood*.

This SOW is provided as an **example of how enterprise AI work is scoped and sequenced**, not as a commitment to build.

2.0 Illustrative AI Models

2.1 Demand & Production Forecasting Model

Purpose Align brewing, packaging, and inventory decisions with real demand signals.

Model Inputs

- Historical sales by SKU, format, and channel
- Distributor depletion and sell-through data
- Seasonal and promotional signals
- Production capacity and lead times

Model Capabilities

- Forecast near-term demand by SKU and market
- Recommend brew and packaging schedules
- Flag overproduction or shortage risk early

Business Impact

- Reduced inventory waste
- Fewer stock-outs in high-performing markets
- Improved cash flow and production planning

Illustrative Timeline: 3–4 months **Illustrative Budget Range:** $75K–$100K

2.2 Distributor & Territory Performance Intelligence Model

Purpose Provide objective visibility into distributor execution and territory performance.

Model Inputs

- Distributor sales and depletion reports
- Account-level performance data
- Market and chain benchmarks
- Sales activity data (where available)

Model Capabilities

- Normalize inconsistent distributor reporting
- Identify underperforming territories or accounts
- Surface high-ROI focus areas

Business Impact

- Stronger distributor accountability
- More effective sales prioritization
- Improved market-level decision-making

Illustrative Timeline: 4–5 months **Illustrative Budget Range:** $90K–$120K

2.3 Product & Channel Mix Optimization Model

Purpose Optimize product mix across draft, packaged, taproom, and retail channels.

Model Inputs

- Sales velocity by SKU and format
- Margin and cost data
- Channel-specific performance
- Seasonal and regional trends

Model Capabilities

- Identify the most profitable product/channel combinations
- Recommend portfolio focus adjustments
- Highlight opportunities for simplification or expansion

Business Impact

- Higher margin per barrel
- Better alignment between brand strategy and production reality
- Reduced operational complexity

Illustrative Timeline: 5–6 months **Illustrative Budget Range:** $100K–$130K

3.0 Deployment Philosophy

Enterprise AI systems are **not built in isolation**.

They are pursued only after:

- Human workflows are mapped and stabilized
- Operational constraints are validated
- Data availability and quality are confirmed

This sequencing prevents automating broken processes and ensures AI investments generate measurable business value.

4.0 Scalability & Value Capture

Once validated, these models may be:

- Operated internally as decision-support systems
- Extended across facilities, brands, or regions
- Adapted for use by other regional breweries with similar profiles

Potential value capture models include:

- Subscription access
- Performance-based fees tied to measurable improvements
- Hybrid licensing and advisory structures

5.0 Closing Note

This document is **illustrative only**.

Its purpose is to demonstrate:

- How enterprise AI work is responsibly scoped
- Why operational clarity must precede automation
- How AI investments are sequenced to reduce risk

The Strategic AI Roadmap and subsequent Architecture Sprint determine **if** and **which** of these models are justified.